

▼ Dataset

WiLI-2018, the Wikipedia language identification benchmark dataset, contains 235000 paragraphs of 235 languages. Each language in this dataset contains 1000 rows/paragraphs. The following is the subset of dataset used:

- English • Arabic • French • Hindi • Urdu • Portuguese • Persian • Pushto • Spanish • Korean • Tamil • Turkish • Estonian • Russian • Romanian • Chinese • Swedish • Latin • Indonesian • Dutch • Japanese • Thai

<https://www.kaggle.com/zarajamshaid/language-identification-datasst>

▼ Importing the libraries

```
import pandas as pd # Pandas
import numpy as np # Numpy
import re # Regular Expressions
import seaborn as sns # Seaborn
import matplotlib.pyplot as plt # Matplotlib
```

+ Code+ Text

▼ Loading the dataset from drive location

```
from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
cd /content/gdrive/MyDrive/NLP/Language Identification

/content/gdrive/MyDrive/NLP/Language Identification
```

▼ Dataset Inspection

```
data = pd.read_csv("dataset.csv")
data.head(10)
```

	Text	language
0	klement gottwaldi surnukeha palsameeriti ning ...	Estonian
1	sebes joseph pereira thomas på eng the jesuit...	Swedish
2	ถนนเจริญกรุง ถนนโรมัน thanon charoen krung l...	Thai
3	விசாகப்பட்டினம் தமிழ்ச்சாங்கத்தை இந்துப் பத்திர...	Tamil
4	de spons behoort tot het geslacht haliclona en...	Dutch
5	エノが行きがかりでバスに乗ってしまい、気分が悪くなった際に助けるが、今すぐバスを降りたいと運...	Japanese
6	tsutinalar ingilizce tsuutina kanadada albert...	Turkish
7	müller mox figura centralis circulorum doctoru...	Latin
8	تمام زیجوبری ذرات کی electric charge برقی بار	Urdu
9	シャーリー・フィールドは、サン・ベルナルド・アベニュー沿い市民センターとrtマーティン高校に...	Japanese

▼ Checking different unique languages in the dataset

```
lang = np.array(pd.unique(data['language']))
print("Total number of langauges:", len(lang))
print("Languages in the dataset:", lang)
```

Total number of langauges: 22
Languages in the dataset: ['Estonian' 'Swedish' 'Thai' 'Tamil' 'Dutch' 'Japanese' 'Turkish' 'Latin' 'Urdu' 'Indonesian' 'Portugese' 'French' 'Chinese' 'Korean' 'Hindi' 'Spanish' 'Pushto' 'Persian' 'Romanian' 'Russian' 'English' 'Arabic']

▼ Checking the count of each language in the dataset

```
data['language'].value_counts()
```

Turkish	1000
Urdu	1000
Swedish	1000
Pushto	1000
Chinese	1000
Romanian	1000
Latin	1000
Estonian	1000
Arabic	1000

Portuguese	1000
Dutch	1000
Hindi	1000
Korean	1000
Japanese	1000
French	1000
Thai	1000
Tamil	1000
Russian	1000
English	1000
Spanish	1000
Persian	1000
Indonesian	1000

Name: language, dtype: int64

▼ Preprocessing the dataset

```
def data_preprocess(text):
    # Removing numbers and symbols
    text = re.sub(r'[!@#$()-_,"%*?;~`0-9]', ' ', text)
    text = re.sub(r'[\[\]]', ' ', text)
    # Lowercasing the text
    text = text.lower()
    return text
```

```
i = 0
# Looping over all the
for text in data['Text']:
    # Updating with the preprocessed dataset
    data['Text'][i] = data_preprocess(text)
    i += 1
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:4: FutureWarning: Possible nested set at position 1
  after removing the cwd from sys.path.
```

▼ Splitting the dataset into train, validation and test sets

```
import math # math library

# Fraction of training examples
train_split = 0.6
# Fraction of testing examples
validation_split = 0.2

# Length of the dataset
length = len(data)

# Splitting the dataset
range1 = math.floor(0.6*length)
range2 = math.floor(0.6*length) + math.floor(0.2*length)
training_data = data[:range1]
validation_data = data[range1: range2]
testing_data = data[range2:]

# Size of train, validation and test dataset
print("Training data length: ", len(training_data))
print("Validation data length: ", len(validation_data))
print("Testing data length: ",len(testing_data))
```

```
Training data length: 13200
Validation data length: 4400
Testing data length: 4400
```

▼ Getting the n-grams from the dataset

The lower and upper boundary of the range of n-values for different word n-grams or char n-grams to be extracted. All values of n such such that $\text{min_n} \leq n \leq \text{max_n}$ will be used. For example an ngram_range of (1, 1) means only unigrams, (1, 2) means unigrams and bigrams, and (2, 2) means only bigrams. n-feat represents number of most important n-gram features to be extracted from corpus.

```
# Import CountVectorizer object
from sklearn.feature_extraction.text import CountVectorizer

# Function to extract n-grams
def get_n_grams(corpus, n_feat, min_n, max_n):
    #fit the n-gram model
    vectorizer = CountVectorizer(analyzer = 'char',
                                ngram_range = (min_n, max_n),
                                max_features = n_feat)

    # Train vectorizer on given corpus
    X = vectorizer.fit_transform(corpus)
    # Get model feature names
    feature_names = vectorizer.get_feature_names()
    return feature_names
```

▼ Building the vocabulary

The following are changeable hyperparameters:

1. `n_feat = 200` # No of n-grams to be extracted from single language
2. `min_n = 3` # min n-gram
3. `max_n = 3` # max n-gram

```
n_feat = 200 # No of n-grams to be extracted from single language
min_n = 3 # min n-gram
max_n = 3 # max n-gram

#obtain n-grams from each language
features_set = set()
lang = pd.unique(training_data['language'])
for l in lang:
    # get corpus filtered by language
    corpus = data[data.language==l]['Text']
    # get n_feat most frequent trigrams
    trigrams = get_n_grams(corpus, n_feat, min_n, max_n)
    # add to set
    features_set.update(trigrams)

#create vocabulary list using feature set
vocab = dict()
for i,f in enumerate(features_set):
    vocab[f]=i
```

▼ Length of vocabulary

```
vocab_length = len(vocab)
print(vocab_length)
```

```
2845
```

▼ Creating feature matrix for training, validation and text data

```
#train count vectoriser using vocabulary
vectorizer = CountVectorizer(analyzer='char',
                             ngram_range=(3, 3),
                             vocabulary=vocab)

#create feature matrix for training set
corpus = training_data['Text']
X = vectorizer.fit_transform(corpus)
feature_names = vectorizer.get_feature_names()
training_feat = pd.DataFrame(data=X.toarray(),columns=feature_names)

#Scale feature matrix
train_min = training_feat.min()
train_max = training_feat.max()
training_feat = (training_feat - train_min)/(train_max-train_min)

#Add target variable
training_feat['language'] = list(training_data['language'])

#create feature matrix for validation set
corpus = validation_data['Text']
X = vectorizer.fit_transform(corpus)

validation_feat = pd.DataFrame(data=X.toarray(),columns=feature_names)
validation_feat = (validation_feat - train_min)/(train_max-train_min)
validation_feat['language'] = list(validation_data['language'])

#create feature matrix for test set
corpus = testing_data['Text']
X = vectorizer.fit_transform(corpus)

testing_feat = pd.DataFrame(data=X.toarray(),columns=feature_names)
testing_feat = (testing_feat - train_min)/(train_max-train_min)
testing_feat['language'] = list(testing_data['language'])
```

```
training_feat.shape
```

```
(13200, 2846)
```

```
validation_feat.shape
```

```
(4400, 2846)
```

```
testing_feat.shape
```

(4400, 2846)

▼ Encode the target vector

```
from sklearn.preprocessing import LabelEncoder
from keras.utils import np_utils

#Fit encoder
encoder = LabelEncoder()
encoder.fit(lang)

def encode(y):
    y_encoded = encoder.transform(y)
    y_dummy = np_utils.to_categorical(y_encoded)
    return y_dummy
```

▼ Building and training the DNN

```
# Importing keras
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout

#Get training data
x = training_feat.drop('language',axis=1)
y = encode(training_feat['language'])
x_val = validation_feat.drop('language',axis=1)
y_val = encode(validation_feat['language'])

#Define model
model = Sequential()
#input dimension is the length of vocabulary
model.add(Dense(512, input_dim=vocab_length, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(256, activation='relu'))
# Output dimension is the length of total number of target languages
model.add(Dense(len(lang), activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

▼ Training the model

Hyperparameters:

1. epochs : No of times dataset is passed forward and backward through the neural network.
2. batch_size : The number of examples in a batch.


```
epochs = 4 # No of times dataset is passed forward and backward through the neural network
batch_size = 100 # The number of examples in a batch

history = model.fit(x, y, epochs=epochs, batch_size=batch_size, validation_data=(x_val, y_val))
```

```
Epoch 1/4
132/132 [=====] - 4s 28ms/step - loss: 0.5276 - accuracy: 0.9148 - val_loss: 0.1027 - val_accuracy: 0.9800
Epoch 2/4
132/132 [=====] - 3s 25ms/step - loss: 0.0715 - accuracy: 0.9849 - val_loss: 0.0783 - val_accuracy: 0.9845
Epoch 3/4
132/132 [=====] - 3s 25ms/step - loss: 0.0303 - accuracy: 0.9933 - val_loss: 0.0863 - val_accuracy: 0.9832
Epoch 4/4
132/132 [=====] - 3s 25ms/step - loss: 0.0146 - accuracy: 0.9973 - val_loss: 0.0881 - val_accuracy: 0.9850
```

▼ Plotting the results

```
# summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



▼ Evaluating on the test set

```
x_test = testing_feat.drop('language',axis=1)
y_test = encode(testing_feat['language'])
history = model.evaluate(x_test, y_test)
```

```
138/138 [=====] - 1s 5ms/step - loss: 0.1034 - accuracy: 0.9786
```

▼ Saving the model in current location

```
model.save('model.h5')
```

▼ Saving the parameters in current location

```
import pickle
filename = 'parameters.sav'
list_of_dumps = [train_max, train_min, vectorizer, feature_names, encoder]
pickle.dump(list_of_dumps, open(filename, 'wb'))
```

Making sample predictions

▼ Detecting the language

```
# Preprocessing each text line
def data_preprocess(text):
    # Removing numbers and symbols
    text = re.sub(r'[!@#$()-_,"%^*?;~`0-9]', ' ', text)
    text = re.sub(r'[\[\]]', ' ', text)
    # Lowercasing the text
    text = text.lower()
    return text

# Function to detect language
def detect_language(text):
    text = data_preprocess(text)
    X = vectorizer.fit_transform([text])
    X_feat = pd.DataFrame(data=X.toarray(),columns=feature_names)
    X_feat = (X_feat - train_min)/(train_max-train_min)
    predicted_my_val = model.predict(X_feat)
    val = np.where(predicted_my_val[0] == np.amax(predicted_my_val[0]))[0]
    print("Detected Language :", encoder.classes_[val[0]])
```

```
my_text = input("Enter the text : ")
my_val = detect_language(my_text)
```

```
Enter the text : De taal is in grammaticaal opzicht verregaand gelijk aan het Hindi dat in Bollywood-films wordt gesproken. De taal die i
Detected Language : Dutch
```